

a-b-testing

January 10, 2026

1 1. Project Overview

1.1 Context

This project focuses on conducting an **A/B Test** for an E-commerce platform. The company has developed a new web page designed to increase the **conversion rate** (the number of users who decide to pay for a product).

1.2 Objectives

The primary goal is to process and analyze interaction logs from both the **Old Page (Control Group)** and the **New Page (Treatment Group)**. The analysis aims to support a data-driven decision-making process regarding the deployment strategy:

- **Deploy:** Launch the new page if it shows a statistically significant improvement.
- **Retain:** Keep the old page if the new version underperforms or shows no significant change.
- **Extend Testing:** Run the experiment longer if the current data is inconclusive to reach a confident decision.

```
[1]: # Setup & Imports
import numpy as np
import pandas as pd
import os

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Statistics
from scipy.stats import chisquare, beta
import statsmodels.api as sm
from statsmodels.stats.proportion import proportions_ztest, proportion_confint
from statsmodels.stats.multipletests import multipletests

# Configuration
pd.set_option('display.max_columns', None)
import warnings
warnings.filterwarnings('ignore')
```

```
print("Libraries imported successfully!")
```

Libraries imported successfully!

2 2. Dataset Description

The analysis utilizes two datasets containing **294,478** user interaction logs.

2.1 2.1. Main Dataset (ab_data.csv)

Feature	Description
user_id	Unique identifier for each user.
timestamp	The timestamp of the user's visit.
group	The experiment group: - control : Users seeing the old_page . - treatment : Users seeing the new_page .
landing_page	The actual web page served to the user (old_page or new_page).
converted	The target variable indicating purchase behavior: - 1: Purchased (Converted). - 0: Did not purchase.

2.2 2.2. Supplemental Data (countries.csv)

This dataset provides location information and can be joined with the main dataset via **user_id**.
* **country**: The country of origin for the user (e.g., US, UK, CA).

```
[2]: main_path = "/kaggle/input/ecommerce-ab-testing-2022-dataset1/
    ↪ecommerce_ab_testing_2022_dataset1/ab_data.csv"
countries_path = "/kaggle/input/ecommerce-ab-testing-2022-dataset1/
    ↪ecommerce_ab_testing_2022_dataset1/countries.csv"

# Load datasets
try:
    df_main = pd.read_csv(main_path)
    df_countries = pd.read_csv(countries_path)

    # print("\nInfo of Main Data")
    # df_main.info()

    # print("\nInfo of Countries Data")
    # df_countries.info()

    print(f"Main Data Shape: {df_main.shape}")
    print(f"Countries Data Shape: {df_countries.shape}")
```

```

print("\nMain Data Head")
display(df_main.head())

print("\nCountries Data Head")
display(df_countries.head())
except FileNotFoundError:
    print("Error: File not found. Please check the paths.")

```

Main Data Shape: (294480, 5)
 Countries Data Shape: (290586, 2)

Main Data Head

	user_id	timestamp	group	landing_page	converted
0	851104	11:48.6	control	old_page	0
1	804228	01:45.2	control	old_page	0
2	661590	55:06.2	treatment	new_page	0
3	853541	28:03.1	treatment	new_page	0
4	864975	52:26.2	control	old_page	1

Countries Data Head

	user_id	country
0	834778	UK
1	928468	US
2	822059	UK
3	711597	UK
4	710616	UK

3 3. Data Engineering Pipeline

To ensure a robust analysis, we construct a data processing pipeline consisting of Data Integration, Deduplication, and Integrity Checks.

3.1 3.1. Data Integration

Goal: Enrich the raw interaction logs with geographical data to enable location-based segmentation later. - **Operation:** Perform a **Left Join** on `user_id`. - **Logic:** We keep all interaction logs from the main dataset and attach country codes where available.

```
[3]: # --- Data Integration ---
# Merge main logs with countries data
df_merged = df_main.merge(df_countries, on="user_id", how="left")

print(f"Shape before merge: {df_main.shape}")
print(f"Shape after merge: {df_merged.shape}")

# Preview data to ensure 'country' column is added
```

```
display(df_merged.head())
```

Shape before merge: (294480, 5)

Shape after merge: (294482, 6)

```
user_id timestamp group landing_page converted country
0 851104 11:48.6 control old_page 0 US
1 804228 01:45.2 control old_page 0 US
2 661590 55:06.2 treatment new_page 0 US
3 853541 28:03.1 treatment new_page 0 US
4 864975 52:26.2 control old_page 1 US
```

3.2 3.2. Temporal Deduplication

Goal: Isolate the user's **first impression** to avoid bias from repeated visits. - **Problem:** Users may visit the landing page multiple times. Including all visits violates the "independence" assumption of statistical tests. - **Strategy:** 1. Keep only the **first** occurrence for each **user_id**. 2. Drop subsequent duplicate entries.

```
[4]: # --- Temporal Deduplication ---
initial_rows = df_merged.shape[0]

# Drop duplicates, keeping the first entry
df_clean = df_merged.drop_duplicates(subset = ['user_id'], keep = 'first')

# Calculate metrics
dropped_rows = initial_rows - df_clean.shape[0]
print(f"Duplicate records removed: {dropped_rows}")
print(f"Unique users remaining: {df_clean.shape[0]}")
```

Duplicate records removed: 3897

Unique users remaining: 290585

3.3 3.3. Mismatch Correction

Goal: Standardize the dataset by aligning `landing_page` with the assigned `group`. - **Assumption:** The `group` assignment (Control/Treatment) is considered the **Source of Truth**. Any mismatch in `landing_page` is treated as a logging error. - **Action:** Instead of dropping valuable user data, we **correct** the `landing_page` values: - If `group == 'control'`, force `landing_page → 'old_page'`. - If `group == 'treatment'`, force `landing_page → 'new_page'`

```
[5]: # --- Mismatch Correction ---
# Identify mismatches before correction
mismatch_mask = (
    ((df_clean['group'] == 'treatment') & (df_clean['landing_page'] == 'old_page')) |
    ((df_clean['group'] == 'control') & (df_clean['landing_page'] == 'new_page')))
```

```

)
print(f"Mismatches detected before correction: {mismatch_mask.sum()}")

df_clean.loc[df_clean['group'] == 'control', 'landing_page'] = 'old_page'
df_clean.loc[df_clean['group'] == 'treatment', 'landing_page'] = 'new_page'

# Verification: Ensure mismatches are gone
recheck_mismatch = (
    ((df_clean['group'] == 'treatment') & (df_clean['landing_page'] ==
    ↪'old_page')) |
    ((df_clean['group'] == 'control') & (df_clean['landing_page'] ==
    ↪'new_page')))
).sum()
print(f"Mismatches remaining after fix: {recheck_mismatch}")

# Create Final Dataframe
df_final = df_clean.copy()

print("=" * 50)
print(f"Final Dataset Shape: {df_final.shape}")
print("=" * 50)

```

```

Mismatched detected before correction: 2044
Mismatched remaining after fix: 0
=====
Final Dataset Shape: (290585, 6)
=====
```

3.4 3.4. Feature Engineering (String Processing)

Goal: Extract session duration from the MM:SS.f string format. - **Strategy:** Split the string directly to extract Minutes and Seconds, avoiding datetime conversion errors. - **Formula:** Time = Minutes × 60 + Seconds.

```
[6]: # Feature Engineering (String Approach)

# Ensure string type
df_final['timestamp'] = df_final['timestamp'].astype(str)

# Split string by ':' to separate Minutes and Seconds
# Format is MM:SS.f -> Part 0 is MM, Part 1 is SS.f
time_parts = df_final['timestamp'].str.split(':', expand=True)

# Convert extracted parts to numeric
minutes = pd.to_numeric(time_parts[0], errors='coerce')
seconds = pd.to_numeric(time_parts[1], errors='coerce')

# Calculate Total Seconds
```

```

df_final['time'] = minutes * 60 + seconds

# Cleanup (Remove rows with bad format that result in NaN)
initial_len = len(df_final)
df_final = df_final.dropna(subset=['time'])

print("-" * 30)
print(f"Rows removed due to bad time format: {initial_len - len(df_final)}")
print(f"Final Dataset Shape: {df_final.shape}")
print("-" * 30)

# Preview
display(df_final[['timestamp', 'time']].head())

```

 Rows removed due to bad time format: 0
 Final Dataset Shape: (290585, 7)

	timestamp	time
0	11:48.6	708.6
1	01:45.2	105.2
2	55:06.2	3306.2
3	28:03.1	1683.1
4	52:26.2	3146.2

4 4. Exploratory Data Analysis (EDA)

We begin by summarizing the key metrics: **Conversion Rate** (Did they buy?) and **Time Spent** (How long did they stay?).

4.1 4.1. Descriptive Statistics

Goal: Reproduce the summary tables to compare performance between the `old_page` and `new_page`. - **n:** Sample size. - **tb (Mean Time):** Average time spent on the page (seconds). - **sd (Std Time):** Standard deviation of time. - **Conversion Rate:** Percentage of users who converted.

[7]: # Descriptive Statistics

```

# General Summary by Landing Page
summary_general = df_final.groupby('landing_page').agg(
    n=('user_id', 'count'),
    conversion_rate=('converted', 'mean'),
    tb_time=('time', 'mean'),
    sd_time=('time', 'std')
).reset_index()

```

```

display(summary_general)

# Detailed Summary by Page & Conversion Status
summary_detailed = df_final.groupby(['landing_page', 'converted']).agg(
    n=('user_id', 'count'),
    tb_time=('time', 'mean'),
    sd_time=('time', 'std')
).reset_index()

display(summary_detailed)

# Quick Observation
p_old = summary_general.loc[summary_general['landing_page'] == 'old_page', 'conversion_rate'].values[0]
p_new = summary_general.loc[summary_general['landing_page'] == 'new_page', 'conversion_rate'].values[0]
print(f"\nObservation: Old Rate ({p_old:.2%}) vs New Rate ({p_new:.2%})")

```

	landing_page	n	conversion_rate	tb_time	sd_time
0	new_page	145353	0.118842	1801.561902	1039.123734
1	old_page	145232	0.120297	1804.922550	1037.313523

	landing_page	converted	n	tb_time	sd_time
0	new_page	0	128079	1801.442659	1039.534724
1	new_page	1	17274	1802.446035	1036.100925
2	old_page	0	127761	1806.096534	1037.039892
3	old_page	1	17471	1796.337502	1039.301750

Observation: Old Rate (12.03%) vs New Rate (11.88%)

4.1.1 Observations from Descriptive Statistics

Based on the summary tables above, we can draw the following preliminary insights:

- Sample Size:** The traffic is evenly distributed (~145,000 users each), confirming valid randomization.
- User Engagement:** Users spent **slightly less time** on the New Page (1801.5s) compared to the Old Page (1804.9s).
- Conversion Performance:** The conversion rate for the New Page (11.88%) is **lower** than the baseline (12.03%).

Preliminary Conclusion: The New Page is underperforming in both engagement and conversion.

4.2 4.2. Visual Analysis

We perform visualization in two stages: 1. **Primary KPI Analysis (Conversion):** Check sample balance, class imbalance, and conversion rates. 2. **Secondary Metric Analysis (Time):** Analyze user engagement depth using Boxplots.

4.2.1 4.2.1. Primary KPI: Conversion Analysis

First, we analyze the Conversion Rate from multiple angles: Overview, Class Distribution, and Confidence Intervals.

```
[8]: # Conversion Analysis

fig, axes = plt.subplots(2, 2, figsize=(16, 10))


sns.countplot(x='landing_page', data=df_final, ax=axes[0, 0], palette="pastel")
axes[0, 0].set_title('Sample Size Distribution (Check Balance)', fontsize=12, fontweight='bold')
for p in axes[0, 0].patches:
    axes[0, 0].annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2., p.get_height()),
                        ha='center', va='bottom', fontsize=10)

sns.histplot(df_final['converted'], kde=False, bins=3, ax=axes[0, 1])
axes[0, 1].set_title('Total Data Distribution (Class Imbalance)', fontsize=12, fontweight='bold')
axes[0, 1].set_xticks([0, 1])
axes[0, 1].set_xticklabels(['Not Converted (0)', 'Converted (1)'])

sns.barplot(x='landing_page', y='converted', data=df_final, ax=axes[1, 0], palette="viridis", errorbar=None)
axes[1, 0].set_title('Conversion Rate (Bar Chart)', fontsize=12, fontweight='bold')
axes[1, 0].set_ylim(0, 0.15)
for p in axes[1, 0].patches:
    axes[1, 0].annotate(f'{p.get_height():.2%}', (p.get_x() + p.get_width() / 2., p.get_height()),
                        ha='center', va='bottom', fontsize=10)

sns.pointplot(x='group', y='converted', data=df_final, ax=axes[1, 1], capsizes=.1, linestyles="--", markers="o", color="b")
axes[1, 1].set_title('Conversion Rate w/ Confidence Intervals', fontsize=12, fontweight='bold')

plt.tight_layout()
plt.show()
```



4.2.2 Observations from Visual Analysis

Based on the dashboard visualizations above, we derive the following key insights regarding the experiment's validity and performance:

1. **Experiment Validity (Sample Balance):** The **Sample Size Distribution** chart confirms that traffic was split evenly between the two groups (~145,000 users each). This validates the randomization process, ensuring that any difference in performance is due to the page design, not traffic volume.
2. **Dataset Characteristics (Class Imbalance):** The **Total Data Distribution** plot highlights a significant class imbalance. The vast majority of users do not convert (Outcome 0), with conversions (Outcome 1) being a relatively rare event (~12%). This suggests that detecting a "lift" requires a highly sensitive test given the large volume of non-events.
3. **Primary KPI (Conversion Performance):**
 - **Bar Chart:** Visually confirms that the `old_page` (12.03%) has a slight edge over the `new_page` (11.88%).
 - **Point Plot (Confidence Intervals):** This is the most critical visual. The downward slope from Control to Treatment indicates a **negative trend**. The confidence intervals (vertical bars) for the two groups do not show a positive overlap, strongly suggesting that the new page is likely performing worse, not better.

Preliminary Conclusion: The visual evidence paints a clear picture: the New Page is **underperforming**. It failed to improve user retention (Time) and appears to negatively impact the primary goal (Conversion Rate). We will now confirm if this observed drop is statistically significant using Hypothesis Testing.

4.2.3 4.2.2. Secondary Metric: Time Analysis (Deep Dive)


Finally, we investigate the **Time Spent** to see if the new page engages users longer, even if they don't buy. - **Boxplot:** To compare medians and detect outliers. - **Histogram:** To understand the density of user engagement.

```
[9]: # Time Analysis
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Chart 1: Boxplot (Distribution & Outliers)
sns.boxplot(x='landing_page', y='time', data=df_final, ax=axes[0],  
            palette="Set2")
axes[0].set_title('Time Spent Distribution (Boxplot)', fontsize=14,  
                  fontweight='bold')
axes[0].set_ylabel('Time (seconds)')

# Chart 2: Histogram (Density)
sns.histplot(data=df_final, x='time', hue='landing_page', element="step",  
             stat="density", common_norm=False, ax=axes[1], palette="Set2")
axes[1].set_title('Time Spent Density (Histogram)', fontsize=14,  
                  fontweight='bold')
axes[1].set_xlabel('Time (seconds)')

plt.tight_layout()
plt.show()
```



4.2.4 Key Insights from Time Analysis

- **No Significant Difference:** Both the Boxplot and Histogram show that the distribution of time spent is identical between the `old_page` and `new_page`.
- **Behavioral Pattern:** The data follows a **uniform distribution**. The median time remains stable at ~ 30 minutes for both groups.
- **Implication:** The redesign failed to alter user engagement behavior.

5 5. Statistical Hypothesis Testing (A/B Testing)

To scientifically answer the question: “Is the new website effectively better than the old one?”, we perform a statistical inference test.

5.1 5.1. Hypothesis Definition

We define the conversion rates for the old and new pages as p_{old} and p_{new} .

- **Null Hypothesis (H_0):** $p_{\text{new}} \leq p_{\text{old}}$ (*The New Page is not better; it is equal to or worse than the Old Page*).
- **Alternative Hypothesis (H_1):** $p_{\text{new}} > p_{\text{old}}$ (*The New Page is strictly better than the Old Page*).
- **Significance Level (α):** 0.05 (95% Confidence Level).

Methodology: Unlike the *Permutation Test* used in smaller datasets, we utilize the **Two-Proportions Z-Test** for this analysis. This is the industry standard for large datasets ($N > 30$) and provides statistically equivalent results to simulation methods but is more computationally efficient.

```
[10]: n_new_converted = df_final[(df_final['landing_page'] == 'new_page') &
    ↵(df_final['converted'] == 1)].shape[0]
n_old_converted = df_final[(df_final['landing_page'] == 'old_page') &
    ↵(df_final['converted'] == 1)].shape[0]

n_new_total = df_final[df_final['landing_page'] == 'new_page'].shape[0]
n_old_total = df_final[df_final['landing_page'] == 'old_page'].shape[0]

counts = np.array([n_new_converted, n_old_converted])
nobs = np.array([n_new_total, n_old_total])

print(f"New Page: {n_new_converted}/{n_new_total} converted (Rate:↵
    ↵{n_new_converted/n_new_total:.4f})")
print(f"Old Page: {n_old_converted}/{n_old_total} converted (Rate:↵
    ↵{n_old_converted/n_old_total:.4f})")
```

New Page: 17274/145353 converted (Rate: 0.1188)

Old Page: 17471/145232 converted (Rate: 0.1203)

5.2 5.2. Test Execution

We utilize the `proportions_ztest` function from the `statsmodels` library to calculate the P-value.

- **Decision Rule:**
 - If $P_{\text{value}} < 0.05$: Reject $H_0 \rightarrow$ The New Page is significantly better.
 - If $P_{\text{value}} \geq 0.05$: Fail to reject $H_0 \rightarrow$ The New Page is not better.

```
[11]: # Run
z_stat, p_value = proportions_ztest(counts, nobs, alternative='larger')
```

```

print(f"Z-Test Results")
print(f"Z-Statistic: {z_stat:.4f}")
print(f"P-Value: {p_value:.4f}")

alpha = 0.05
print("\nSTATISTICAL DECISION")

if p_value < alpha:
    print("Result: REJECT H0.")
    print("Interpretation: The New Page is statistically better than the Old\u20acPage.")
else:
    print("Result: FAIL TO REJECT H0.")
    print(f"Interpretation: The difference is not statistically significant,\u20ac(p-value {p_value:.4f} > {alpha}).")
    print("Evidence suggests the New Page is NOT better.")

```

```

Z-Test Results
Z-Statistic: -1.2091
P-Value: 0.8867

STATISTICAL DECISION
Result: FAIL TO REJECT H0.
Interpretation: The difference is not statistically significant (p-value 0.8867
> 0.05).
Evidence suggests the New Page is NOT better.

```

5.3 5.3. Final Business Conclusion

Statistical Analysis: The test resulted in a p-value of approximately 0.8867, which is significantly higher than our alpha level of 0.05. * Mathematically, this means we **fail to reject the Null Hypothesis (H_0)**. * The probability that the New Page is “better” by random chance is close to zero, given that the observed data actually shows a decline.

Business Recommendation: 1. **Performance:** The New Page conversion rate (11.88%) is statistically equivalent to (or slightly worse than) the Old Page (12.03%). 2. **Engagement:** As seen in the EDA section, the new design failed to increase user session duration. 3. **Action:** The company **should NOT deploy** the new landing page. * **Reasoning:** Implementing the new page would incur deployment costs without providing any revenue uplift.